

# MA8701 Advanced methods in statistical inference and learning

## L1: Introduction and core concepts

Mette Langaas

1/9/23

### Table of contents

<b>1</b>	<b>Course philosophy</b>	<b>2</b>
<b>2</b>	<b>Learning</b>	<b>4</b>
2.1	Learning outcome . . . . .	4
<b>3</b>	<b>Course topics</b>	<b>4</b>
3.1	Part 1: Core concepts [3 weeks] . . . . .	6
3.2	Part 2: Shrinkage and regularization [3 weeks] . . . . .	6
3.3	Part 3: Ensembles [4 weeks] . . . . .	6
3.4	Part 4: XAI [2 weeks] . . . . .	7
3.5	Part 5: Closing [2 weeks] . . . . .	7
3.6	Some observations about the course . . . . .	7
3.7	“Required” previous knowledge . . . . .	7
3.8	Course elements . . . . .	8
3.9	Class activity . . . . .	8
3.10	Learning methods and activities . . . . .	9
<b>4</b>	<b>Tentative plan for part 1</b>	<b>9</b>
4.1	L1 . . . . .	9
4.2	L2 . . . . .	9
4.3	W2 . . . . .	10
4.4	W3 . . . . .	10
<b>5</b>	<b>Core concepts</b>	<b>10</b>
5.1	Notation . . . . .	10
5.2	Random variables and random vectors . . . . .	11

5.3	Training set . . . . .	11
5.4	Validation and test data . . . . .	11
5.5	Group discussion . . . . .	12
5.6	Regression and MLR . . . . .	12
<b>6</b>	<b>Statistical decision theoretic framework</b>	<b>12</b>
6.1	Squared error loss . . . . .	13
6.1.1	What if the joint distribution is multivariate normal? . . . . .	14
6.1.2	Approximate linear model . . . . .	14
6.1.3	kNN and conditional expectation . . . . .	14
6.2	Absolute loss . . . . .	14
<b>7</b>	<b>Discussion and conclusions</b>	<b>15</b>
<b>8</b>	<b>Exercises</b>	<b>15</b>
8.0.1	1: Law of total expectation and total variance . . . . .	15
8.0.2	2: Quadratic loss and decision theoretic framework . . . . .	15
8.0.3	3: Curse of dimensionality . . . . .	15
8.0.4	4: Key results from MLR . . . . .	16
<b>9</b>	<b>Solutions to exercises</b>	<b>17</b>
9.0.1	1: Law of total E and Var . . . . .	17
9.0.2	2: Quadratic loss . . . . .	17
9.0.3	3: Curse of dimensionality . . . . .	17
9.0.4	4: Key results from MLR . . . . .	17
<b>10</b>	<b>Reference links</b>	<b>18</b>
<b>11</b>	<b>Bibliography</b>	<b>18</b>

Course homepage: <https://wiki.math.ntnu.no/ma8701/2023v/start>

---


## 1 Course philosophy

“Important new algorithms often arise outside the world of professional statisticians: neural nets, support vector machines, and boosting are three famous examples. None of this is surprising. New sources of data, satellite imagery for example, or medical microarrays, inspire novel methodology from the observing scientists. The early literature tends toward the enthusiastic, with claims of enormous applicability and power.”

In the second phase, statisticians try to locate the new methodology within the framework of statistical theory. In other words, they carry out the statistical inference part of the cycle, placing the new methodology within the known Bayesian and frequentist limits of performance. This is a healthy chain of events, good both for the hybrid vigor of the statistics profession and for the further progress of algorithmic technology.”

Efron and Hastie (2016) (page 10).

---



## 2 Learning

### 2.1 Learning outcome

#### 1. Knowledge

- Understand and explain the central theoretical aspects in statistical inference and learning.
- Understand and explain how to use methods from statistical inference and learning to perform a sound data analysis.
- Be able to evaluate strengths and weaknesses for the methods and choose between different methods in a given data analysis situation.

[Take home message in MA8701 in 2021 according to the students](#)

---

#### 2. Skills

Be able to analyse a dataset using methods from statistical inference and learning in practice (using R or Python), and give a good presentation and discussion of the choices done and the results found.

#### 3. Competence

- The students will be able to participate in scientific discussions, read research presented in statistical journals.
  - They will be able to participate in applied projects, and analyse data using methods from statistical inference and learning.
- 

## 3 Course topics


The starting point is that we cover important parts of

The Elements of Statistical Learning: Data Mining, Inference, and Prediction, Second Edition (Springer Series in Statistics, 2009) by Trevor Hastie, Robert Tibshirani, and Jerome Friedman.

but, since the book is from 2009 (2nd edition, 12th corrected printing from 2017) this means that for many topic we need (to be up to date) additional selected material in the form of book chapters and research articles.

Download book at <https://hastie.su.domains/ElemStatLearn/download.html> (this is 12th corrected printing) or sign in at NTNU vpn at [Springer](#) (but I am not sure if this is an earlier print without the version 12 corrections).

---



---

### 3.1 Part 1: Core concepts [3 weeks]

Sort out assumed background knowledge, and learn something new

- Notation
- Repetition of core concepts (regression and classification)
- Statistical decision theoretic framework (partly new) ESL 2.4
- Model selection and model assessment - including bias-variance trade-off (mostly new) ESL 7
- Handbook of Missing Data Methodology (parts of Chapters 11-12, partly 13) and Flexible Imputation of Missing Data (parts of Chapters 2-4)

---

### 3.2 Part 2: Shrinkage and regularization [3 weeks]

or “Regularized linear and generalized linear models”, with focus on the ridge and lasso regression (in detail).

- ESL 3.2.3, 3.4, 3.8, 4.4.4.
- Hastie, Tibshirani, Wainwright (HTW): “Statistical Learning with Sparsity: The Lasso and Generalizations”. Selected sections from Chapters 1, 2, 3, 4, 6.
- Selective inference (articles)

---

### 3.3 Part 3: Ensembles [4 weeks]

- trees, bagging and random forests
- xgboost
- general ensembles (including super learner)
- hyper-parameter tuning

Selected Chapters in ESL (8.7, 8.8, 9.2, parts of 10, 15, 16) and several articles.

---

### 3.4 Part 4: XAI [2 weeks]

Lectured by Kjersti Aas <https://www.nr.no/~kjersti/>.

Interpretable Machine Learning: A Guide for Making Black Box Models Explainable, Molnar (2019), with the following topics:

- LIME,
- partial dependence plots,
- Shapley values,
- relative weights and
- counterfactuals.

### 3.5 Part 5: Closing [2 weeks]

---

### 3.6 Some observations about the course

- Mainly a frequentist course, but some of the concepts and methods have a Bayesian version that might give insight into why and how the methods work, then Bayesian methods will be used.
  - Focus is on regression and classification, and unsupervised learning is not planned to be part of the course.
  - The required previous knowledge is listed because this is a PhD-course designed for statistics students. The background make the students go past an overview level of understanding of the course parts (move from algorithmic to deep understanding).
- 

### 3.7 “Required” previous knowledge

- TMA4267 Linear statistical methods
- TMA4268 Statistical learning
- TMA4295 Statistical inference
- TMA4300 Computer intensive statistical methods
- TMA4315 Generalized linear models
- Good understanding and experience with R, or with Python, for statistical data analysis.
- Knowledge of markdown for writing reports and presentations (Rmarkdown/Quarto, Jupyter).
- Skills in group work - possibly using git or other collaborative tools.

---

### 3.8 Course elements

Course wiki at <https://wiki.math.ntnu.no/ma8701/2023v/start>

- Lectures
- Problem sets to work on between lectures.
- Office hours and/or [mattelab.math.ntnu.no](https://mattelab.math.ntnu.no/)?
- Study techniques (share)
- Ethical considerations
- Compulsory work
- Final individual oral exam in May

The learning material is also available at <https://github.com/mettelang/MA8701V2023>.

---

### 3.9 Class activity

Aim: get to know each other - to improve on subsequent group work!

```
while (at least one student not presented)
  lecturer give two alternatives, you choose one.
  lecturer choose a few students to present their view
  together with giving their name and study programme
  (and say if they are looking for group members)
```

- 
- Dog person or cat person?
  - When performing logistic regression - do you then say you do statistical learning or machine learning?
  - I will show you the result of a descriptive analysis: table summary or graphical display?
  - Learning something new: read a book or watch a video?
  - Analysing data: R or python?
  - Analysing data: report p-values and or confidence intervals
  - In class: taking notes or not?



- Use camel case or snake case for programming?

camel: writing compound words such that each word in the middle of the phrase begins with a capital letter, with no intervening spaces or punctuation. “camelCase” or “CamelCase”.

snake: writing compound words where the elements are separated with one underscore character (\_\_) and no spaces, with each element’s initial letter usually lower cased within the compound and the first letter either upper- or lower case as in “foo\_bar”

---

### 3.10 Learning methods and activities

Herbert A. Simon (Cognitive science, Nobel Laureate): *Learning results from what the student does and thinks and only from what the student does and thinks. The teacher can advance learning only by influencing what the student does to learn.*

---

## 4 Tentative plan for part 1

(progress may be faster or slower than indicated)

### 4.1 L1

Notation, regression and statistical theoretic framework

- Notation (ESL Ch 2.2)
- Regression - should not be new (ESL Ch 3, except 3.2.3, 3.2.4, 3.4, 3.7, 3.8)
- Statistical decision theoretic framework for regression (ESL 2.4)

### 4.2 L2

Continue with the same framework but for classification, if time also bias-variance trade-off

- Classification - should not be new (ESL Ch 4.1-4.5, except 4.4.4)
  - Statistical decision theoretic framework for classification (ESL 2.4)
  - and the bias-variance trade-off
-

## 4.3 W2

L3-4: Then, cover new aspects for

- Model selection and assessment (ESL Ch 7.1-7.6, 7.10-7.12)

## 4.4 W3

L5-6

- How to handle missing data in data analyses
- 

# 5 Core concepts

## 5.1 Notation

(mainly from ESL 2.2)

We will only consider supervised methods.

- Response  $Y$  (or  $G$ ): dependent variable, outcome, usually univariate (but may be multivariate)
  - quantitative  $Y$ : for regression
  - qualitative, categorical  $G$ : for classification, some times dummy variable coding used (named one-hot coding in machine learning)
- Covariates  $X_1, X_2, \dots, X_p$ : “independent variables”, predictors, features
  - continuous, discrete: used directly
  - categorical, discrete: often dummy variable coding used

We aim to construct a rule, function, learner:  $f(X)$ , to predict  $Y$  (or  $G$ ).

---

Random variables and (column) vectors are written as uppercase letters  $X$ , and  $Y$ , while observed values are written with lowercase  $(x, y)$ . (Dimensions specified if needed.)

Matrices are presented with uppercase boldface:  $X$ , often  $N \times (p + 1)$ .

ESL uses boldface also for  $x_j$  being a vector of all  $N$  observations of variable  $j$ , but in general vectors are not boldface and the vector of observed variables for observation  $i$  is just  $x_i$ .

---

## 5.2 Random variables and random vectors

Both the response *and covariates* will be considered to be random, and drawn from some joint distribution  $P(X_1, X_2, \dots, X_p, Y) = P(X, Y)$  or  $P(X, G)$ .

Joint to conditional and marginal distribution:  $P(X, Y) = P(Y | X)P(X)$  or  $P(Y | X = x)P(X = x)$  or

$$P(Y = y, X = x) = P(Y = y | X = x)P(X = x)$$

Maybe brush up on this?

### Resources

- From TMA4268: [Module 2 - Random vectors](#)
- From TMA4267: [Part 1: Multivariate random variables and the multivariate normal distribution](#)

---

## 5.3 Training set

(ESL 2.1)

A set of size  $N$  of independent pairs of observations  $(x_i, y_i)$  is called the *training set* and often denoted  $\mathcal{T}$ . Here  $x_i$  may be a vector.

The training data is used to estimate the unknown function  $f$ .

## 5.4 Validation and test data

Validation data is used for *model selection* (finding the best model among a candidate set).

Test data is used for *model assessment* (assess the performance of the fitted model on future data).

We will consider theoretical results, and also look at different ways to split or resample available data.

More in ESL Chapter 7.

## 5.5 Group discussion

Two core regression methods are multiple linear regression (MLR) and  $k$ -nearest neighbour (kNN).

For the two methods

- Set up the formal definition for  $f$ , and model assumptions made
  - What top results do you remember? Write them down.
  - What are challenges?
- 

## 5.6 Regression and MLR

See also the exercises!

### Resources

(mostly what we learned in TMA4267, or ESL Ch 3, except 3.2.3, 3.2.4, 3.4, 3.7, 3.8)

- From TMA4268: [Overview](#) and in particular [Module 3: Linear regression](#)
- From TMA4315: [Overview](#) and in particular [Module 2: MLR](#)

For  $k$ NN see also Problem 1 of the [TMA4268 2018 exam](#) with [solutions](#)

---

## 6 Statistical decision theoretic framework

(ESL Ch 2.4, regression part)

is a mathematical framework for developing models  $f$  - and assessing optimality.

First, regression:

- $X \in \mathfrak{R}^p$
- $Y \in \mathfrak{R}$
- $P(X, Y)$  joint distribution of covariates and responses

Aim: find a function  $f(X)$  for predicting  $Y$  from some inputs  $X$ .

Ingredients: Loss function  $L(Y, f(X))$  - for *penalizing errors in the prediction*.

Criterion for choosing  $f$ : Expected prediction error (EPE)

---


$$\text{EPE}(f) = \mathbb{E}_{X,Y}[L(Y, f(X))] = \int_{x,y} L(y, f(x))p(x, y)dxdy$$

Choose  $f$  to minimize the  $\text{EPE}(f)$ .

---

Q: Why do we not involve the distribution of the random variable  $f(X)$ , but instead the distribution of  $X$ ?

Law of the unconscious statistician (from our introductory course in statistics): [https://en.wikipedia.org/wiki/Law\\_of\\_the\\_unconscious\\_statistician](https://en.wikipedia.org/wiki/Law_of_the_unconscious_statistician) and [Thematics pages TMA4240/45](#).

What is the most popular loss function for regression?

---

## 6.1 Squared error loss

$$\text{EPE}(f) = \mathbb{E}_{X,Y}[L(Y, f(X))] = \mathbb{E}_X \mathbb{E}_{Y|X}[(Y - f(X))^2 | X]$$

We want to minimize EPE, and see that it is sufficient to minimize  $\mathbb{E}_{Y|X}[(Y - f(X))^2 | X]$  for each  $X = x$  (pointwise):

$$f(x) = \operatorname{argmin}_c \mathbb{E}_{Y|X}[(Y - c)^2 | X = x]$$

This gives as result the conditional expectation - the best prediction at any point  $X = x$ :

$$f(x) = \mathbb{E}[Y | X = x]$$

Proof: by differentiating and setting equal 0.

But, do we know this conditional distribution? In practice: need to estimate  $f$ .

---

### 6.1.1 What if the joint distribution is multivariate normal?

Conditionally (known from before):

if we assume that  $(X, Y) \sim N_{p+1}(\mu, \Sigma)$  then we have seen (TMA4267) that  $E(Y | X)$  is linear in  $X$  and  $\text{Cov}(Y | X)$  is independent of  $X$ .

[Properties of the mvN](#)

---

### 6.1.2 Approximate linear model

But, also if we assume an approximate linear model:  $f(x) \approx x^T \beta$

Marginally:  $\text{argmin}_{\beta} E[(Y - X^T \beta)^2]$  gives  $\beta = E[XX^T]^{-1} E[XY]$  (now random vectors).

We may replace expectations with averages in training data to estimate  $\beta$ .

This is not conditional on  $X$ , but we have assumed a linear relationship.

---

### 6.1.3 kNN and conditional expectation

Local conditional mean for observations in  $T$  close to  $\mathbf{x}_0$ :

$$\hat{f}(\mathbf{x}_0) = \frac{1}{k} \sum_{i \in N_k(\mathbf{x}_0)} Y_i$$

---

## 6.2 Absolute loss

Regression with absolute (L1) loss:  $L(Y, f(X)) = |Y - f(X)|$  gives  $\hat{f}(x) = \text{median}(Y | X = x)$ .

Proof: for example pages 8-11 of [https://getd.libs.uga.edu/pdfs/ma\\_james\\_c\\_201412\\_ms.pdf](https://getd.libs.uga.edu/pdfs/ma_james_c_201412_ms.pdf)

## 7 Discussion and conclusions

- What are key take home messages from today's teaching session?
  - What do you plan to do before the next teaching session?
  - Feedback on today's teaching session?
- 

## 8 Exercises

### 8.0.1 1: Law of total expectation and total variance

This is to get a feeling of the joint and conditional distributions, so that we understand expected value notation with joint, conditional and marginal distributions.

Give a derivation of the law of total expectation:

$$E[X] = E[E(X | Y)]$$

and the law of total variance:

$$\text{Var}[X] = E\text{Var}[X | Y] + \text{Var}E[X | Y]$$

(There is also a law of total covariance.)

### 8.0.2 2: Quadratic loss and decision theoretic framework

Show that  $f(x) = E[Y | X = x]$  for the quadratic loss.

---

### 8.0.3 3: Curse of dimensionality

Read pages 22-23 and then answer Exercise 2.3 - which is to "Derive equation (2.24)."

Important take home messages:

- All sample points are close to an edge of the sample.
- If data are uniformly distributed in an hypercube in  $p$  dimensions, we need to cover  $r^{1/p}$  of the the range of each input variable to capture a fraction  $r$  of the observations.

#### 8.0.4 4: Key results from MLR

(These results are known from TMA4267 and TMA4315, but useful to brush up on?)

Assume we have a data set with independent observation pairs  $(y_i, \mathbf{x}_i)$  for  $i = 1, \dots, N$ .

$$\mathbf{Y} = X\beta +$$

where  $\mathbf{Y}$  is a  $N \times 1$  random column vector,  $\mathbf{X}$  a  $N \times (p + 1)$  design matrix with row for observations  $(\mathbf{x}_i^T)$  and columns for covariates, and  $\varepsilon$  a  $N \times 1$  random column vector

The assumptions for the classical linear model is:

1.  $E(\varepsilon) = \mathbf{0}$ .
2.  $\text{Cov}(\varepsilon) = E(\varepsilon\varepsilon^T) = \sigma^2\mathbf{I}$ .
3. The design matrix has full rank,  $\text{rank}(\mathbf{X}) = (p + 1)$ .

The classical *normal* linear regression model is obtained if additionally

4.  $\varepsilon \sim N_n(\mathbf{0}, \sigma^2\mathbf{I})$  holds.

For random covariates these assumptions are to be understood conditionally on  $\mathbf{X}$ .

##### 8.0.4.1 a) Regression parameter estimator

Derive the least squares (or the maximum likelihood estimator) for  $\hat{\beta}$ .

##### 8.0.4.2 b) Properties of regression parameter estimator

Derive the distribution of  $\hat{\beta}$  (when assumption 4 is true).

##### 8.0.4.3 c) Estimator for variance

Explain how we may find the restricted maximum likelihood estimator for  $\sigma^2$ . Which distribution is used for inference for  $\sigma^2$ ?

##### 8.0.4.4 d) Hypothesis test

How would you test the hypothesis  $H_0 : \beta_j = 0$  against  $H_1 : \beta_j \neq 0$ ?



#### 8.0.4.5 e) Explanability

Explanability is now very important - but then we usually talk about black box models. How would you explain the impact of each covariate in a multiple linear regression model? Can you give the proportion of the variability that each variable is responsible for explaining?

---

## 9 Solutions to exercises

Please try yourself first, or take a small peek - and try some more - before fully reading the solutions. Report errors or improvements to [Mette.Langaas@ntnu.no](mailto:Mette.Langaas@ntnu.no). (The solutions given here are very similar to the UiO STK-IN4300 solutions, see link under References.)

### 9.0.1 1: Law of total E and Var

[Try first yourself](#)

### 9.0.2 2: Quadratic loss

Page 8 of [https://getd.libs.uga.edu/pdfs/ma\\_james\\_c\\_201412\\_ms.pdf](https://getd.libs.uga.edu/pdfs/ma_james_c_201412_ms.pdf) or [Mettes notes](#)

### 9.0.3 3: Curse of dimensionality

[2.3](#)

### 9.0.4 4: Key results from MLR

#### 9.0.4.1 a) Regression parameter estimator

Both methods are written out in [these class notes from TMA4267/8](#). More on likelihood-version here: [TMA4315 GLM Module 2](#).

#### 9.0.4.2 b) Properties of regression estimator

[Page 3 of classnotes from TMA4268](#)

#### 9.0.4.3 c) Estimator for variance

[TMA4315 GLM Module 2](#)

#### 9.0.4.4 d) Hypothesis test

[Page 4 of classnotes from TMA4268](#)

#### 9.0.4.5 e)

Standardized regression coefficients, estimated regression coefficients divided by their standard errors, is the most used solution. But, also popular is the decomposition of the  $R^2$  - easy for orthogonal design matrix, but not easy for correlated covariates. The LMG-method of Grömping (2007) (decomposing  $R^2$ ) gives a solution that is also valid with correlated covariates, that is identical to the Shapley value of Part 4 (more later).

## 10 Reference links

- [ESL official errata](#): and choose “Errata” in the left menu
- [ESL notes and solutions to exercises](#)
- [ESL solutions from UiO](#)
- [Camel\\_case](#)
- [Snake\\_case](#)

## 11 Bibliography

Efron, Bradley, and Trevor Hastie. 2016. *Computer Age Statistical Inference - Algorithms, Evidence, and Data Science*. Cambridge University Press. <https://hastie.su.domains/CA SI/>.

Grömping, U. 2007. “Estimators of Relative Importance in Linear Regression Based on Variance Decomposition.” *The American Statistician* 61: 139–47.

Hastie, Trevor, Robert Tibshirani, and Jerome Friedman. 2009. *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*. Vol. 2. Springer series in statistics New York. [hastie.su.domains/ElemStatLearn](https://hastie.su.domains/ElemStatLearn).

Molnar, Christoph. 2019. *Interpretable Machine Learning: A Guide for Making Black Box Models Explainable*.