

MA8701 Advanced methods in statistical inference and learning

Part 3: Ensembles. L15: Stacked ensembles

Mette Langaas

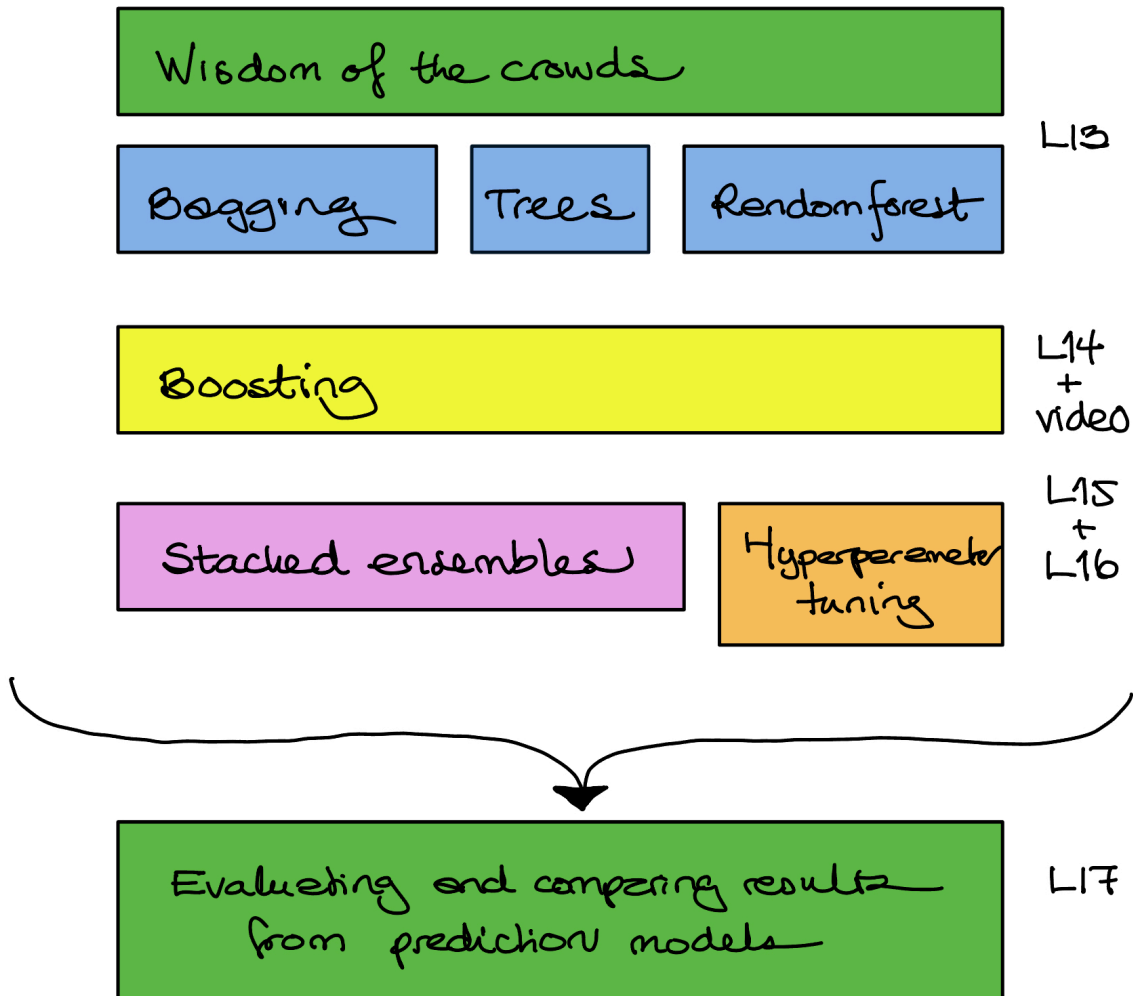
3/5/23

Table of contents

1	Before we start	2
1.1	Literature	2
1.2	Supporting literature	3
2	Ensembles - overview	3
3	Stacked ensembles	3
3.1	What is it?	4
3.2	Development:	4
3.3	Ingredients:	4
3.4	Algorithm	5
3.5	The metalearning	6
4	Examples	6
4.1	Simulation examples	6
4.2	Real data	7
5	Theoretical result	7
5.1	Uncertainty in the ensemble	7
5.2	Other issues	8
6	R example from Superlearner package	8
7	R example from H2o-package	13
7.1	Now the default metalearner	17

Course homepage: <https://wiki.math.ntnu.no/ma8701/2023v/start>

1 Before we start



1.1 Literature

- Erin Le Dell (2015): Scalable Ensemble Learning and Computationally Efficient Variance Estimation. PhD Thesis, University of California, Berkeley. or <https://github.com/led>

1.2 Supporting literature

- Breiman (1996)
 - Laan, Polley, and Hubbard (2007)
 - Polley, Rose, and Laan (2011)
-

2 Ensembles - overview

(ELS Ch 16.1)

With ensembles we want to build *one prediction model* which combines the strength of *a collection of models*.

These models may be simple base models - or more elaborate models.

We have studied bagging - where we use the bootstrap to repeatedly fit a statistical model, and then take a simple average of the predictions (or majority vote). Here the base models can be trees - or other type of models.

Random forest is a version of bagging with trees, with trees made to be different (decorrelated).

We have studied boosting, where the models are trained on sequentially different data - from residuals or gradients of loss functions - and the ensemble members cast weighted votes (down-weighted by a learning rate). We have observed that there are many hyperparameters that need to be tuned to optimize performance.

3 Stacked ensembles

aka super learner or generalized stacking

3.1 What is it?

The Stacked Ensembles is an algorithm that combines

- multiple, (typically) diverse prediction methods (learning algorithms) called *base learners* (first-level) into a
 - a second-level *metalearner* - which can be seen as a *single* method.
-

3.2 Development:

- 1992: stacking introduced for neural nets by Wolpert
 - 1996: adapted to regression problems by Breiman - but only for one type of methods at once (CART with different number of terminal nodes, GLMs with subset selection, ridge regression with different ridge penalty parameters) Breiman (1996)
 - 2006: proven to have asymptotic theoretical oracle property by Laan, Polley, and Hubbard (2007)
 - 2015: extensions in phd thesis by Erin LeDell LeDell (2015)
-

3.3 Ingredients:

- *Training data* (level-zero data) $O_i = (X_i, Y_i)$ of N i.i.d observations.
 - A total of L *base learning algorithms* Ψ^l for $l = 1, \dots, L$, each from some algorithmic class and each with a specific set of model parameters.
 - A *metalearner* Φ is used to find an *optimal combination* of the L base learners.
-

3.4 Algorithm

Step 1: Produce level-one data Z

- a) Divide the training data X randomly into V roughly-equally sized validation folds $X_{(1)}, \dots, X_{(V)}$. V is often 5 or 10. (The responses Y are also needed.)
- b) For each base learner Ψ^l perform V -fold cross-validation to produce prediction.

This gives the level-one data set Z consisting prediction of all the level-zero data - that is a matrix with N rows and L columns.

What could the base learners be?

“Any” method that produces a prediction - “all” types of problems.

- linear regression
- lasso
- cart
- random forest with mtry=value 1
- random forest with mtry=value 2
- xgboost with hyperparameter set 1
- xgboost with hyperparameter set 2
- neural net with hyperparameter set 1

Step 2: Fit the metalearner

- a) The starting point is the level-one prediction data Z together with the responses (Y_1, \dots, Y_N) .
- b) The metalearner is used to estimate the weights given to each base learner: $\hat{\eta}_i = \alpha_1 z_{1i} + \dots + \alpha_L z_{Li}$.

What could the metalearner be?

-
- the mean (bagging)
 - constructed by minimizing the
 - squared loss (ordinary least squares) or

- non-negative least squares (most popular)
- ridge or lasso regression
- logistic regression (for binary classification)
- constructed by minimizing 1-ROC-AUC

(Class notes: Study Figure 3.2 from Polley, Rose, and Laan (2011) and/or Figure 1 from Laan, Polley, and Hubbard (2007))

3.5 The metalearning

Some observations

- The term *discrete super learner* is used if the base learner with the lowest risk (i.e. CV-error) is selected.
 - Since the predictions from multiple base learners may be highly correlated - the chosen method should perform well in that case (i.e. ridge and lasso).
 - When minimizing the squared loss it has been found that adding a non-negativity constraint $\alpha_l \leq 0$ works well,
 - and also the additivity constraint $\sum_{l=1}^L \alpha_l = 1$ - the ensemble is a *convex combination* of the base learners.
 - Non-linear optimization methods may be employed for the metalearner if no existing algorithm is available
 - Historically a regularized linear model has “mostly” been used
 - For classification the logistic response function can be used on the linear combination of base learners (Figure 3.2 Polley, Rose, and Laan (2011)).
-

4 Examples

4.1 Simulation examples

(Class notes: Study Figure 3.3 and Table 3.2 from Polley, Rose, and Laan (2011))

4.2 Real data

(Class notes: Study Figure 3.4 and Table 3.3 from P@Polley2011. RE=MSE relative to the linear model OLS.)

5 Theoretical result

LeDell (2015) (page 6)

- Oracle selector: the estimator among all possible weighted combinations of the base prediction function that minimizes the risk under the *true data generating distribution*.
 - The *oracle result* was established for the Super Learner by Laan, Polley, and Hubbard (2007)
 - If the *true prediction function* cannot be represented by a combination of the base learners (available), then “optimal” will be the closest linear combination that would be optimal if the true data-generating function was known.
 - The oracle result require an *uniformly bounded loss function*. Using the convex restriction (sum alphas =1) implies that if each based learner is bounded so is the convex combination. In practice: truncation of the predicted values to the range of the outcome in the training set is sufficient to allow for unbounded loss functions
-

5.1 Uncertainty in the ensemble

(Class notes: Study “Road map” 2 from Polley, Rose, and Laan (2011))

- Add an outer (external) cross validation loop (where the super learner loop is inside). Suggestion: use 20-fold, especially when small sample size.
 - Overfitting? Check if the super learner does as well or better than any of the base learners in the ensemble.
 - Results using *influence functions* for estimation of the variance for the Super Learner are based on asymptotic variances in the use of V -fold cross-validation (see Ch 5.3 of LeDell (2015))
-

5.2 Other issues

- Many different implementations available, and much work on parallel processing and speed and memory efficient execution.
 - Super Learner implicitly can handle hyperparameter tuning by including the same base learner with different model parameter sets in the ensemble.
 - Speed and memory improvements for large data sets involves subsampling, and the R `subsemble` package is one solution, the H2o package another.
-

6 R example from Superlearner package

Comment - this package is still in use, but the `h2o-superlearner` might be more “easy” to use.

Code is copied from [Guide to SuperLearner](#) and the presentation follows this guide. The data used is the Boston housing dataset from MASS, but with the median value of a house dichotomized into a classification problem.

Observe that only 150 of the 560 observations is used (to speed up things, but of course that gives less accurate results).

```
data(Boston, package = "MASS")
#colSums(is.na(Boston)) # no missing values
outcome = Boston$medv
# Create a dataframe to contain our explanatory variables.
data = subset(Boston, select = -medv)
#Set a seed for reproducibility in this random sampling.
set.seed(1)
# Reduce to a dataset of 150 observations to speed up model fitting.
train_obs = sample(nrow(data), 150)
# X is our training sample.
x_train = data[train_obs, ]
# Create a holdout set for evaluating model performance.
# Note: cross-validation is even better than a single holdout sample.
x_holdout = data[-train_obs, ]
# Create a binary outcome variable: towns in which median home value is > 22,000.
outcome_bin = as.numeric(outcome > 22)
y_train = outcome_bin[train_obs]
y_holdout = outcome_bin[-train_obs]
```



```
table(y_train, useNA = "ifany")
```

```
y_train
 0  1
92 58
```

Then checking out the possible functions and how they differ from their “original versions”.

```
listWrappers()
```

```
[1] "SL.bartMachine"      "SL.bayesglm"        "SL.biglasso"
[4] "SL.caret"           "SL.caret.rpart"     "SL.cforest"
[7] "SL.earth"           "SL.extraTrees"      "SL.gam"
[10] "SL.gbm"             "SL.glm"             "SL.glm.interaction"
[13] "SL.glmnet"          "SL.ipredbagg"        "SL.kernelKnn"
[16] "SL.knn"             "SL.ksvm"            "SL.lda"
[19] "SL.leekasso"        "SL.lm"              "SL.loess"
[22] "SL.logreg"          "SL.mean"            "SL.nnet"
[25] "SL.nnls"            "SL.polymars"        "SL.qda"
[28] "SL.randomForest"    "SL.ranger"          "SL.ridge"
[31] "SL.rpart"           "SL.rpartPrune"      "SL.speedglm"
[34] "SL.speedlm"         "SL.step"            "SL.step.forward"
[37] "SL.step.interaction" "SL.stepAIC"         "SL.svm"
[40] "SL.template"        "SL.xgboost"

[1] "All"
[1] "screen.corP"         "screen.corRank"      "screen.glmnet"
[4] "screen.randomForest" "screen.SIS"          "screen.template"
[7] "screen.ttest"        "write.screen.template"
```

```
# how does SL.glm differ from glm? obsWeight added to easy use the traning fold in the CV
SL.glm
```

```
function (Y, X, newX, family, obsWeights, model = TRUE, ...)
{
  if (is.matrix(X)) {
    X = as.data.frame(X)
  }
  fit.glm <- glm(Y ~ ., data = X, family = family, weights = obsWeights,
```

```

        model = model)
    if (is.matrix(newX)) {
        newX = as.data.frame(newX)
    }
    pred <- predict(fit.glm, newdata = newX, type = "response")
    fit <- list(object = fit.glm)
    class(fit) <- "SL.glm"
    out <- list(pred = pred, fit = fit)
    return(out)
}
<bytecode: 0x14c2979e0>
<environment: namespace:SuperLearner>

# min and not 1sd used, again obsWeights, make sure model matrix correctly specified
SL.glmnet

function (Y, X, newX, family, obsWeights, id, alpha = 1, nfolds = 10,
        nlambdas = 100, useMin = TRUE, loss = "deviance", ...)
{
    .SL.require("glmnet")
    if (!is.matrix(X)) {
        X <- model.matrix(~-1 + ., X)
        newX <- model.matrix(~-1 + ., newX)
    }
    fitCV <- glmnet::cv.glmnet(x = X, y = Y, weights = obsWeights,
        lambda = NULL, type.measure = loss, nfolds = nfolds,
        family = family$family, alpha = alpha, nlambdas = nlambdas,
        ...)
    pred <- predict(fitCV, newx = newX, type = "response", s = ifelse(useMin,
        "lambda.min", "lambda.1se"))
    fit <- list(object = fitCV, useMin = useMin)
    class(fit) <- "SL.glmnet"
    out <- list(pred = pred, fit = fit)
    return(out)
}
<bytecode: 0x14c2f1350>
<environment: namespace:SuperLearner>

```

The fitting lasso to check what is being done. The default metalearner is “method.NNLS” (both for regression and two-class classification - probably then for linear predictor NNLS?).

```
set.seed(1)
sl_lasso=SuperLearner(Y=y_train, X=x_train,family=binomial(),SL.library="SL.glmnet")
sl_lasso
```

Call:

```
SuperLearner(Y = y_train, X = x_train, family = binomial(), SL.library = "SL.glmnet")
```

	Risk	Coef
SL.glmnet_All	0.08484849	1

```
#str(sl_lasso)
sl_lasso$cvRisk
```

```
SL.glmnet_All
0.08484849
```

Now use lasso and randomforest, and also add the average of ys just as the benchmark.

```
set.seed(1)
sl=SuperLearner(Y=y_train, X=x_train,family=binomial(),SL.library=c("SL.mean","SL.glmnet",
sl
```

Call:

```
SuperLearner(Y = y_train, X = x_train, family = binomial(), SL.library = c("SL.mean",
"SL.glmnet", "SL.randomForest"))
```

	Risk	Coef
SL.mean_All	0.23773937	0.000000
SL.glmnet_All	0.08725786	0.134252
SL.randomForest_All	0.07213058	0.865748

```
sl$times$everything
```

user	system	elapsed
1.489	0.029	1.521

Our ensemble give weight 0.13 to lasso and 0.86 to the random forest. (The guide used a different implementation of the random forest called ranger, and got 0.02 and 0.98.)

Predict on the part of the dataset not used for the training.

```
pred=predict(sl,x_holdout=x_holdout,onlySL=TRUE)
str(pred)
```

List of 2

```
$ pred          : num [1:150, 1] 0.3029 0.07 0.97847 0.00726 0.00523 ...
..- attr(*, "dimnames")=List of 2
.. ..$ : chr [1:150] "505" "324" "167" "129" ...
.. ..$ : NULL
$ library.predict: num [1:150, 1:3] 0.387 0.387 0.387 0.387 0.387 ...
..- attr(*, "dimnames")=List of 2
.. ..$ : chr [1:150] "505" "324" "167" "129" ...
.. ..$ : chr [1:3] "SL.mean_All" "SL.glmnet_All" "SL.randomForest_All"
```

```
summary(pred$pred)
```

```
      V1
Min.   :0.0003034
1st Qu.:0.0183955
Median :0.1135270
Mean   :0.3855066
3rd Qu.:0.9036164
Max.   :0.9955802
```

```
summary(pred$library.predict)
```

SL.mean_All	SL.glmnet_All	SL.randomForest_All
Min. :0.3867	Min. :0.0000014	Min. :0.0000
1st Qu.:0.3867	1st Qu.:0.0244935	1st Qu.:0.0160
Median :0.3867	Median :0.2063204	Median :0.1020
Mean :0.3867	Mean :0.3866667	Mean :0.3853
3rd Qu.:0.3867	3rd Qu.:0.8169726	3rd Qu.:0.9123
Max. :0.3867	Max. :0.9997871	Max. :0.9980

Add now an external cross-validation loop - only using the training data. Here the default $V = 10$ is used for the inner loop, and we set the value for the outer loop (here $V = 3$ for speed).

```
system.time({cv_sl=CV.SuperLearner(Y=y_train, X=x_train,V=10,family=binomial(),SL.library=
```

```
      user  system elapsed
15.003    0.202   15.207
```

```
summary(cv_sl)
```

Call:

```
CV.SuperLearner(Y = y_train, X = x_train, V = 10, family = binomial(), SL.library = c("SL.me
  "SL.glmnet", "SL.randomForest"))
```

Risk is based on: Mean Squared Error

All risk estimates are based on $V = 10$

	Algorithm	Ave	se	Min	Max
	Super Learner	0.077041	0.0123412	0.0223396	0.12156
	Discrete SL	0.079783	0.0132389	0.0245396	0.12151
	SL.mean_All	0.242535	0.0093204	0.1947874	0.29619
	SL.glmnet_All	0.088109	0.0152056	0.0098891	0.14402
	SL.randomForest_All	0.073960	0.0115466	0.0245396	0.12151

See the guide for more information on running multiple versions of one base learner, and parallelisation.

7 R example from H2o-package

<https://docs.h2o.ai/h2o/latest-stable/h2o-docs/data-science/stacked-ensembles.html>

Python examples available from the same page

The Higgs boson data is used - but which version is not specified, maybe this <https://archive.ics.uci.edu/ml/datasets/HIGGS> or a specifically made data set. The problem is binary, so maybe to detect signal vs noise.

Default metalearner: Options include 'AUTO' (GLM with non negative weights; if validation_frame is present, a lambda search is performed)

```
h2o.init()
```

Connection successful!

R is connected to the H2O cluster:

```
H2O cluster uptime:      1 days 3 hours
H2O cluster timezone:    Europe/Oslo
H2O data parsing timezone: UTC
H2O cluster version:     3.40.0.1
H2O cluster version age: 25 days
H2O cluster name:        H2O_started_from_R_mettela_bze126
H2O cluster total nodes: 1
H2O cluster total memory: 2.96 GB
H2O cluster total cores: 10
H2O cluster allowed cores: 10
H2O cluster healthy:     TRUE
H2O Connection ip:       localhost
H2O Connection port:     54321
H2O Connection proxy:    NA
H2O Internal Security:   FALSE
R Version:               R version 4.2.2 (2022-10-31)
```

```
# Import a sample binary outcome train/test set into H2O
train <- h2o.importFile("https://s3.amazonaws.com/erin-data/higgs/higgs_train_10k.csv")
```

```
|
|
|
|====| 0%
|
|=====| 5%
|
|=====| 20%
|
|=====| 32%
|
|=====| 50%
|
|=====| 67%
```

```
|
|=====| 87%
|
|=====| 100%
```

```
test <- h2o.importFile("https://s3.amazonaws.com/erin-data/higgs/higgs_test_5k.csv")
```

```
|
|
| 0%
|=====| 25%
|
|=====| 72%
|
|=====| 100%
```

```
# Identify predictors and response
y <- "response"
x <- setdiff(names(train), y)

# For binary classification, response should be a factor
train[, y] <- as.factor(train[, y])
test[, y] <- as.factor(test[, y])

print(dim(train))
```

```
[1] 10000    29
```

```
print(colnames(train))
```

```
[1] "response" "x1"      "x2"      "x3"      "x4"      "x5"
[7] "x6"       "x7"      "x8"      "x9"      "x10"     "x11"
[13] "x12"      "x13"     "x14"     "x15"     "x16"     "x17"
[19] "x18"      "x19"     "x20"     "x21"     "x22"     "x23"
[25] "x24"      "x25"     "x26"     "x27"     "x28"
```

```
print(dim(test))
```

```
[1] 5000    29
```

```
# Number of CV folds (to generate level-one data for stacking)
n folds <- 5
```

```
# There are a few ways to assemble a list of models to stack together:
# 1. Train individual models and put them in a list
```

```
# 1. Generate a 2-model ensemble (GBM + RF)
```

```
# Train & Cross-validate a GBM
```

```
my_gbm <- h2o.gbm(x = x,
                  y = y,
                  training_frame = train,
                  distribution = "bernoulli",
                  ntrees = 10,
                  max_depth = 3,
                  min_rows = 2,
                  learn_rate = 0.2,
                  n folds = n folds,
                  keep_cross_validation_predictions = TRUE,
                  seed = 1)
```

```
|
|
|
|=====| 67%
|
|=====| 100%
```

```
# Train & Cross-validate a RF
```

```
my_rf <- h2o.randomForest(x = x,
                          y = y,
                          training_frame = train,
                          ntrees = 50,
                          n folds = n folds,
```



```
keep_cross_validation_predictions = TRUE,
seed = 1)
```

```
|
|
|
|=====| 0%
|
|=====| 9%
|
|=====| 26%
|
|=====| 42%
|
|=====| 58%
|
|=====| 74%
|
|=====| 90%
|
|=====| 100%
```

7.1 Now the default metalearner

AUTO: glm with non-negative weights

```
# Train a stacked ensemble using the GBM and RF above
ensemble <- h2o.stackedEnsemble(x = x,
                                y = y,
                                training_frame = train,
                                base_models = list(my_gbm, my_rf))
```

```
|
|
|
|=====| 0%
|
|=====| 100%
```

```

# default metalearner_transform should be NONE
#print(summary(ensemble))
#ensemble@model
# Eval ensemble performance on a test set
perf <- h2o.performance(ensemble, newdata = test)

# Compare to base learner performance on the test set
perf_gbm_test <- h2o.performance(my_gbm, newdata = test)
perf_rf_test <- h2o.performance(my_rf, newdata = test)
baselearner_best_auc_test <- max(h2o.auc(perf_gbm_test), h2o.auc(perf_rf_test))
ensemble_auc_test <- h2o.auc(perf)
print(sprintf("Best Base-learner Test AUC:  %s", baselearner_best_auc_test))

```

```
[1] "Best Base-learner Test AUC:  0.769204725074508"
```

```
print(sprintf("Ensemble Test AUC:  %s", ensemble_auc_test))
```

```
[1] "Ensemble Test AUC:  0.773144298176816"
```

```

# [1] "Best Base-learner Test AUC:  0.76979821502548"
# [1] "Ensemble Test AUC:  0.773501212640419"

# Generate predictions on a test set (if neccessary)
pred <- h2o.predict(ensemble, newdata = test)

```

```

|
|
|
|=====| 100%

```

```
print(head(pred))
```

	predict	p0	p1
1	0	0.6839178	0.3160822
2	1	0.5825428	0.4174572
3	1	0.5869431	0.4130569

```

4      1 0.1927212 0.8072788
5      1 0.4509384 0.5490616
6      1 0.3275080 0.6724920

```

```
metalearner_model
```

```
Model Details:
```

```
=====
```

```
H2OBinomialModel: glm
```

```
Model ID: metalearner_AUTO_StackedEnsemble_model_R_1677945156774_1824
```

```
GLM Model: summary
```

```

      family link                                regularization number_of_predictors_total
1 binomial logit Elastic Net (alpha = 0.5, lambda = 8.399E-5 )                                2
                                     training_frame
1 levelone_training_StackedEnsemble_model_R_1677945156774_1824

```

```
Coefficients: glm coefficients
```

```

              names coefficients standardized_coefficients
1              Intercept      -3.603549              0.149102
2 GBM_model_R_1677945156774_1086      3.298011              0.493334
3 DRF_model_R_1677945156774_1214      3.809905              0.701246

```

```
# Train a stacked ensemble using the GBM and RF above
```

```

ensemble <- h2o.stackedEnsemble(x = x,
                                y = y,
                                training_frame = train,
                                base_models = list(my_gbm, my_rf),
                                metalearner_transform = "Logit")

```

```

|
|
|
|=====| 100%

```

```
#print(summary(ensemble))
```

```
#print(ensemble@model)
```

```
# Eval ensemble performance on a test set
```

```
perf <- h2o.performance(ensemble, newdata = test)
```

```
# Compare to base learner performance on the test set
perf_gbm_test <- h2o.performance(my_gbm, newdata = test)
perf_rf_test <- h2o.performance(my_rf, newdata = test)
baselearner_best_auc_test <- max(h2o.auc(perf_gbm_test), h2o.auc(perf_rf_test))
ensemble_auc_test <- h2o.auc(perf)
print(sprintf("Best Base-learner Test AUC:  %s", baselearner_best_auc_test))
```

```
[1] "Best Base-learner Test AUC:  0.769204725074508"
```

```
print(sprintf("Ensemble Test AUC:  %s", ensemble_auc_test))
```

```
[1] "Ensemble Test AUC:  0.773096033881535"
```

```
# Generate predictions on a test set (if neccessary)
pred <- h2o.predict(ensemble, newdata = test)
```

```
|
|
|
|=====| 100%
```

```
print(head(pred))
```

	predict	p0	p1
1	0	0.6739209	0.3260791
2	1	0.5814741	0.4185259
3	1	0.5826643	0.4173357
4	1	0.1971804	0.8028196
5	1	0.4561659	0.5438341
6	1	0.3365841	0.6634159

```
$metalearner_model
Model Details:
=====
```

```

H2OBinomialModel: glm
Model ID:  metalearner_AUTO_StackedEnsemble_model_R_1677945156774_1830
GLM Model: summary
      family link                                regularization number_of_predictors_total
1 binomial logit Elastic Net (alpha = 0.5, lambda = 3.885E-4 )                                2
                                     training_frame
1 levelone_training_StackedEnsemble_model_R_1677945156774_1830

Coefficients: glm coefficients
              names coefficients standardized_coefficients
1              Intercept      -0.053725                0.154528
2 GBM_model_R_1677945156774_1086      0.791767                0.515081
3 DRF_model_R_1677945156774_1214      0.845217                0.731991

```

8 References

- Breiman, Leo. 1996. “Stacked Regressions.” *Machine Learning* 24 (1): 49–64. <https://doi.org/10.1007/BF00117832>.
- Hastie, Trevor, Robert Tibshirani, and Jerome Friedman. 2009. *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*. Vol. 2. Springer series in statistics New York. hastie.su.domains/ElemStatLearn.
- Laan, Mark J. van der, Eric C Polley, and Alan E. Hubbard. 2007. *Statistical Applications in Genetics and Molecular Biology* 6 (1). <https://doi.org/doi:10.2202/1544-6115.1309>.
- LeDell, Erin. 2015. “Scalable Ensemble Learning and Computationally Efficient Variance Estimation.”
- Polley, Eric C., Sherri Rose, and Mark J. van der Laan. 2011. “Super Learning.” In *Targeted Learning: Causal Inference for Observational and Experimental Data*, 43–66. New York, NY: Springer New York. https://doi.org/10.1007/978-1-4419-9782-1_3.