

MA8701 Advanced methods in statistical inference and learning

Part 3: Ensembles. L16: Hyperparameter tuning

Mette Langaas

3/8/23

Table of contents

1	Before we start	1
1.1	Literature	1
2	Choosing hyperparameters	2
2.1	Surrogate methods	4
2.2	Bayesian optimization	5
2.2.1	Gaussian processes	5
2.2.2	Multivariate normal distribution	7
2.2.3	Acquisition function: Expected improvement	7
2.2.4	Algorithm for Bayesian optimization of a function f	9
2.3	Extension	10
2.3.1	Example	10
2.3.2	Suggested software	16
2.4	Design of experiments and response surface methodology	16
3	References	16

Course homepage: <https://wiki.math.ntnu.no/ma8701/2023v/start>

1 Before we start

1.1 Literature

- Hyperparameter tuning with Bayesian Optimization. Frazier (2018): “A tutorial on Bayesian optimization”, <https://arxiv.org/abs/1807.02811>: Sections 1,2,3,4.1, 5: only

the section “Noisy evaluations”, 6,7.

- G. A. Lujan-Moreno, P. R. Howard, O. G. Rojas and D. C. Montgomery (2018): Design of experiments and response surface methodology to tune machine learning hyperparameters, with a random forest case- study. Expert Systems with Applications. 109, 195-205.

2 Choosing hyperparameters

- What are *hyperparameters*?
- Which hyperparameters have we encountered in the course so far?

Hyperparameters are parameters than cannot be directly estimated from data. This may be model parameters (the λ in lasso) or parameters that influence the fitting of the model (e.g. related to some optimization algorithm).

We have already studied hyperparameter tuning for the lasso, ridge, elastic net, random forest, and boosting - with the use of cross-validation of some loss function for a predefined set (grid) of hyperparameter values.

The use of ensembles like the stacked ensemble (Super Learner) may be seen as an alternative to hyperparameter tuning.

Overview of advices from Berent Å.S. Lunde on tuning parameters in xgboost (from 2021):

Ways to speed-up computation:

- Higher learning-rate, then tune the number of boosting iterations.
- When tuning, do not use too high k in k-fold CV (for both speed and also to avoid high variance), or drop CV and use a validation set.
- Speedups with histogram algorithms of order n (avoid exact enumeration of all splits, $n \log n$).
- Use sparsity when possible!

Comments on hyperparameters:

- Learning rate (eta in xgb): Set as low as computation times allow. Typically in between 0.01 and 0.1 is sufficient.
- Number of trees (boosting iterations): Very much affected by the learning rate. Most important parameter to tune, given a learning rate. Maximum depth of trees: start low, then increase. High values takes significantly longer to train. Think about the problem, and the number of possible interaction effects. If max-depth = J, then interactions among J-1 features is possible. How much is needed?
- Gamma: Very traditional tree-hyperparameter to tune. Tune.
- Colsampling for tree is usually sufficient.
- subsample: in 0.5-0.9
- Min child weight: Something I usually do not think about. Default=1 (low) works. For me

Every problem is different, and there will exist exceptions to the above guidelines. *I look forward to a world without hyperparameters.* (Berent Å.S. Lunde)

The hyperparameters may be continuous (penalty parameter), discrete (number of layers in a neural network, applying early stopping or not) or categorical (choose between different optimizers).

The choice of hyperparameters is important, and will often directly affect the model complexity, and unwise choices may lead to overfitting.

Hyperparameter tuning is performed using a separate validation set or by cross-validation. Different loss functions or selection criteria may be used (MSE, ROC-AUC, misclassification rate, ...).

The hyperparameter tuning is often referred to as a black-box optimization because we (usually) only calculate loss function values (with CV) and do not get to compute gradients.

What may be challenges with hyperparameter optimization?

Some challenges with hyperparameter optimization (Feurer and Hutter, Ch1):

- expensive evaluation of the model under study (large networks, large data sets)
- unclear which of possibly many hyperparameters that need to be selected carefully (refer to the discussion for xgboost)

- gradient of selection criterion with respect to the hyperparameters not (generally) available, and criterion not convex or smooth in the hyperparameters
 - and the need for external validation or CV
-

There exist many ways to *group* methods for hyperparameter tuning. One way to look at this is (Kuhn and Silge, 2021, Ch 12)

- grid search: specify a set of possible values a priori and investigate only these values, choose the value where the chosen selection criterion is optimal. This is also called “model free methods”.
- iterative search: start with a set of values, fit/evaluate some (surrogate) model (might also be the loss function), and based on this choose new values to evaluate next.

For grid search also methods for *speeding up calculations* exists - for example by stopping evaluation at a grid point where the loss is seen to be high after some CV-folds, for example the method of *racing* described by Kuhn and Silge, Ch 13.4.

(Class notes: see example from Kuhn and Silge, 2021, Spacefilling grid vs global search)

2.1 Surrogate methods

We will look at two types of surrogate models: Bayesian regression with Gaussian processes (in Bayesian optimization) and regression-type models in response surface methods.

2.2 Bayesian optimization

Bayesian optimization is an iterative method - where we start with evaluating some loss function at some predefined set of points in the hyperparameter space. New position in the hyperparameter space are chosen iteratively.

Two key ingredients:

- a surrogate model (we will only look at Bayesian regression with Gaussian processes) to fit to the observed values of the loss function in the hyperparameter space
 - an *acquisition* function to decide a new point in the hyperparameter space to evaluate next
-

Underlying idea: given some “observations” in the hyperparameter space, the task is to decide where to place a new point. We should try a point where:

- we expect a good value and/or
- we have little information so far

To do that we need information on both expected value *and* variance - or preferably the distribution of the loss function for our problem.

We now look at the multivariate Gaussian distribution and conditional distribution, a Gaussian process

2.2.1 Gaussian processes

(Eidsvik 2017, page 6-7, note in TMA4265)

A Gaussian process is defined for

- times or locations x_i , $i = 1, \dots, n$ in \mathfrak{R}^d , where
- $Y_i = Y(x_i)$ is a random variable at x_i
- such that $Y = (Y_1, \dots, Y_n)$ is multivariate Gaussian.

The process is *first order (mean) stationary* if $E(Y(x)) = \mu$ for all x , and this can be extended to depend on covariates.

The process is *second order stationary* if $\text{Var}(Y(x)) = \sigma^2$ for all x and the correlation $\text{Corr}(Y(x), Y(x'))$ only depends on differences between x and x' .

The multivariate Gaussian distribution is defined by the mean and covariance alone.

2.2.1.1 Correlation functions

(Eidsvik 2017, page 7, Frazier 2018, Ch 3.1)

Correlation functions are also referred to as *kernels*.

We assume that points at positions close to each other have a stronger correlation than point far apart.

Power exponential or Gaussian kernel

$$\text{Corr}(Y(x), Y(x')) = \exp(-\phi_G \|x - x'\|^2)$$

where the L2 distance is used and ϕ_G is a parameter that determine the decay in the correlations.

Matern-type kernel

$$\text{Corr}(Y(x), Y(x')) = (1 + \phi_M \|x - x'\|) \exp(-\phi_M \|x - x'\|)$$

now with decay-describing parameter ϕ_M .

The parameters of the kernels need to be estimated, see Ch 3.2 of Frazier 2018 (who use a slightly different parameterization). We will just assume that these parameters are known.

(Class notes: study Figure 4 and 5 of Eidsvik, 2018.)

2.2.1.2 From correlations into covariance matrix

For simplicity assume that $d = 1$. The number of positions to consider is n .

To get from correlation function to a $n \times n$ covariance matrix first construct a $n \times n$ matrix of distances for each pair of positions, denote this H .

For the Matern-type correlation function the covariance matrix can then be written

$$\Sigma = \sigma^2 (1 + \phi_M H) \otimes \exp(-\phi_M H)$$

where \otimes is elementwise multiplication.

See Eidsvik (2018, Ch 3.2 and 3.3) for how to build covariance matrices in an efficient way.

2.2.2 Multivariate normal distribution

aka multivariate Gaussian distribution. Known from TMA4265 and TMA4267.

The random vector $Y_{p \times 1}$ is multivariate normal N_p with mean μ and (positive definite) co-variate matrix Σ . The pdf is:

$$f(Y) = \frac{1}{(2\pi)^{\frac{p}{2}} |\Sigma|^{\frac{1}{2}}} \exp\left\{-\frac{1}{2}(Y - \mu)^T \Sigma^{-1} (Y - \mu)\right\}$$

Six useful properties of the mvN - we need number 6.

Let $Y_{(p \times 1)}$ be a random vector from $N_p(\mu, \Sigma)$.

1. The graphical contours of the mvN are ellipsoids (can be shown using spectral decomposition).
2. Linear combinations of components of Y are (multivariate) normal (can be easily proven using moment generating functions MGF).
3. All subsets of the components of Y are (multivariate) normal (special case of the above).

-
4. Zero covariance implies that the corresponding components are independently distributed (can be proven using MGF).
 5. $A\Sigma B^T = 0 \Leftrightarrow AY$ and BY are independent.
 6. The conditional distributions of the components are (multivariate) normal.

$$Y_2 \mid (Y_1 = Y_1) \sim N_{p2}(\mu_2 + \Sigma_{21}\Sigma_{11}^{-1}(Y_1 - \mu_1), \Sigma_{22} - \Sigma_{21}\Sigma_{11}^{-1}\Sigma_{12}).$$

2.2.3 Acquisition function: Expected improvement

(Frazier 2018 page 7)

Thought experiment:

- 1) we have evaluated our function at all possible points x , and must return a solution based on what we already have evaluated. If the evaluation is noise-less we need to return the point with the largest observed value f .

- 2) Correction: We may perform one more evaluation. If we choose x we observe $f(x)$, and the best point before that was f_n^* . The improvement at the new observation is then

$$\max(f(x) - f_n^*, 0)$$

(In class study Figure 1 of Frazier 2018)

-
- 3) We define the *expected improvement* as

$$\text{EI}_n(x) = \mathbb{E}_n[\max(f(x) - f_n^*, 0)]$$

where the expectation is taken at the posterior distribution given that we have evaluated f at n observations x_1, \dots, x_n , and the posterior distribution is that f conditional on $x_1, \dots, x_n, y_1, \dots, y_n$ is normal with mean $\mu_n(x)$ and variance $\sigma_n^2(x)$.

-
- 4) How to evaluate the expected improvement? Integration by parts gives

$$\begin{aligned} \text{EI}_n(x) = & \max(\mu_n(x) - f_n^*, 0) + \sigma_n(x) \phi\left(\frac{\mu_n(x) - f_n^*}{\sigma_n(x)}\right) \\ & - \text{abs}(\mu_n(x) - f_n^*) \Phi\left(\frac{\mu_n(x) - f_n^*}{\sigma_n(x)}\right) \end{aligned}$$

$\mu_n(x) - f_n^*$ is expected proposed vs previously best

- 5) We choose to evaluate the point with the largest expected improvement

$$x_{n+1} = \text{argmax} \text{EI}_n(x)$$

Is often found using quasi-Newton optimization.

2.2.4 Algorithm for Bayesian optimization of a function f

(Frazier 2018, page 3, noise-free evaluation)

Place a Gaussian process prior on f .

Observe f at n_0 points from some experimental design. Set $n = n_0$.

while $n \leq N$ **do**

Update the posterior on f with all available data

Let x_n be a maximizer of the acquisition function over x , computed using the current posterior

Observe $y_n = f(x_n)$

Increment n

end while

Return a solution: a point with largest $f(x)$ or the point with the largest posterior mean

What does the steps mean?

- Gaussian prior: choose (estimate?) mean and correlation function for the problem.
- Observe n_0 points: calculate the loss function at each of the points (remark: we have noise)
- Update the posterior: calculate the conditional distribution for f for a new point given the observed loss at all previously observed points
- Acquisition function: find $\text{argmaxEI}_n(x)$.

(Class notes: Figure 1 of Frazier 2018.)

-
- For a point x we model the distribution of $f(x)$,
 - which is normally distributed with mean $\mu_n(x)$ and variance $\sigma_n^2(x)$. The mean and variance is found from the conditional distribution.
 - With 95% credibility interval $\mu_n(x) \pm 1.95\sigma_n(x)$.
 - The width of the credibility interval at observations is 0.
-

2.3 Extension

What is the objection function is not observed noise-less?


Independent normal error term ε can be added to the previously defined $Y = f(x)$ to make a new $Y = f(x) + \varepsilon$. This (only) adds a diagonal term to the covariance matrix, and it is common to assume that the variance is the same for all x and treat the variance as a hyperparameter.

2.3.1 Example

(Kuhn and Silge, Ch 14, the example is for SVM)

First just grid search to test what is best value for `mtry`

```
data(Boston, package = "MASS")
# first using a grid
tune_grid <- expand.grid(
  mtry = (1:13))
# ntree=seq(100,500,length=10)) # how to also include ntree? primary only mtry, how to de
tune_control <- caret::trainControl(
  method = "oob", # cross-validation #eller cv
  #number = 3, # with n folds
  verboseIter = FALSE, # no training log
  allowParallel = FALSE # FALSE for reproducible results
)
rf_tune <- caret::train(
  medv~crim+zn+indus+chas+nox+rm+age+dis+rad+tax+ptratio+black+lstat,
  data=Boston,
  na.action=na.roughfix,
  trControl = tune_control,
  tuneGrid = tune_grid,
  method = "rf", # rf is randomForest, checked at #vhttp://topepo.github.io/caret/train-mo
  verbose = TRUE
)
tuneplot <- function(x, probs = .90) {
  ggplot(x) +
    coord_cartesian(ylim = c(quantile(x$results$RMSE, probs = probs), min(x$results$RMSE))
    theme_bw()
}
tuneplot(rf_tune)
```



```
rf_tune$bestTune
```

```
mtry
7    7
```

The R the function `tune_bayes` is available in the package `tune`, and requires that the analyses is done with a workflow. Default in the GP is exponential correlation function, but first we try the Matern.

```
tree_rec <- recipe(medv~crim+zn+indus+chas+nox+rm+age+dis+rad+tax+ptratio+black+lstat, data)

tune_spec <- rand_forest( # parsnip interface to random forests models
  mode="regression",
  mtry = tune(),
  trees = tune(),
  # min_n = tune()
) %>%
# set_mode("regression") %>%
# set_engine("ranger",objective="reg:rmse") # errors with ranger
```

```

set_engine("randomForest") # randomforest ok

tune_wf <- workflow() %>%
  add_recipe(tree_rec) %>%
  add_model(tune_spec)

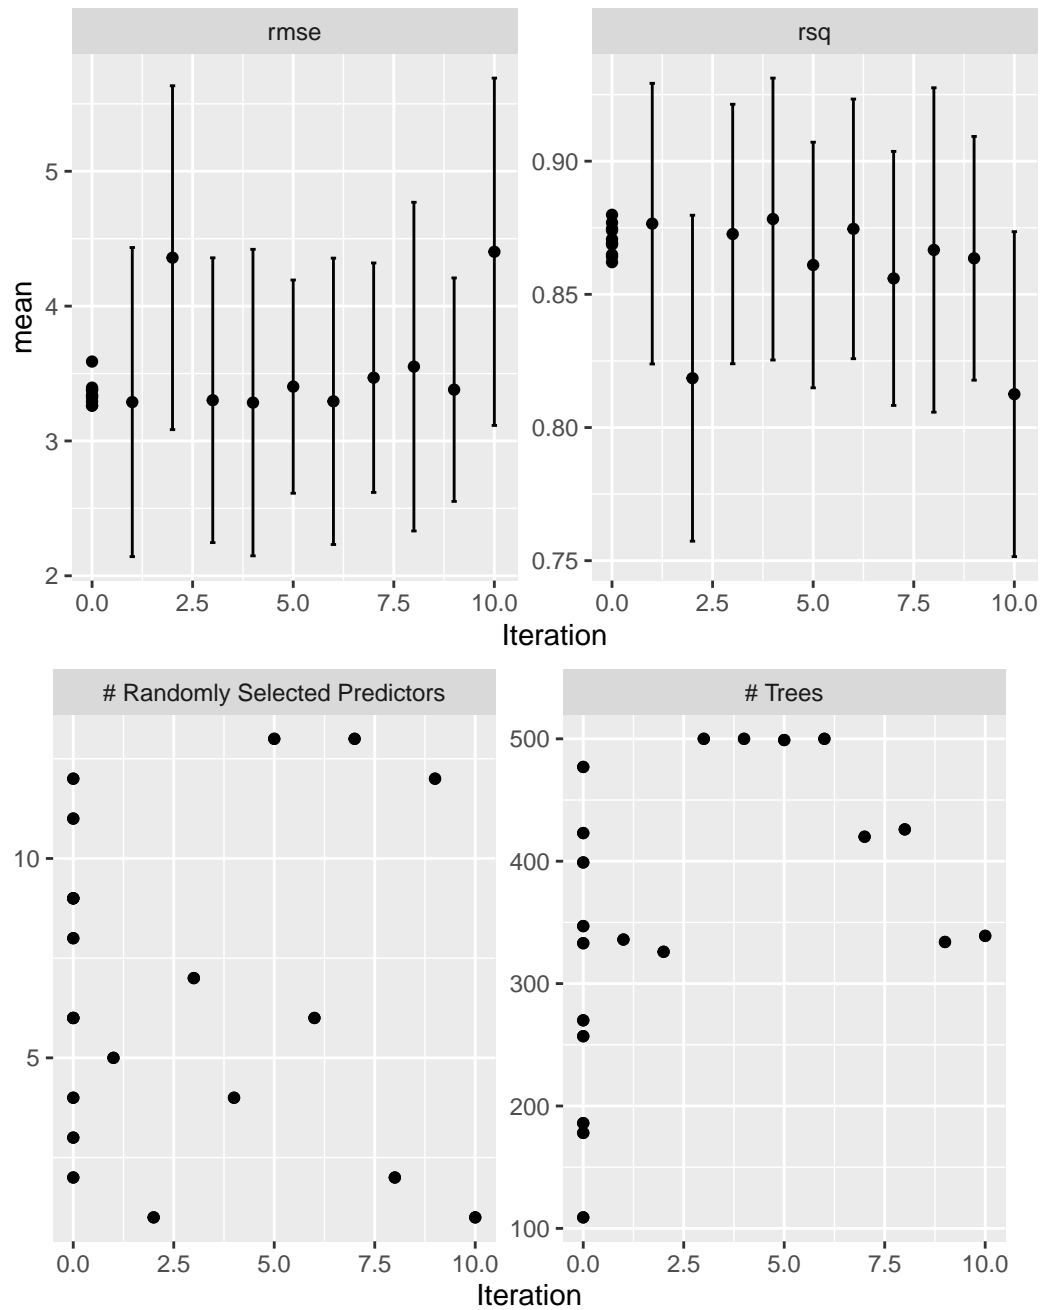
tune_param <- tune_spec%>%
  parameters%>%
  update(mtry=mtry(c(1L,13L)),trees=trees(c(100L,500L)))

vfold <- vfold_cv(Boston, v = 5)
# then trying BO
ctrl <- control_bayes(verbose = TRUE)
bayesres<- tune_bayes(tune_wf,
  resamples = vfold,
  #metrics = rmse,
  corr=list(type="matern",nu=5/2),
  #default in corr_mat(GPfit) is "exponential" power 1.95
  initial = 10,
  param_info = tune_param,
  iter = 10,
  objective=exp_improve(),
  control = ctrl
)
dput(bayesres,"bayesres.dd")

```

A tibble: 10 x 9

	mtry	trees	.metric	.estimator	mean	n	std_err	.config	.iter
	<int>	<int>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>	<int>
1	4	333	rmse	standard	3.26	5	0.439	Preprocessor1_Model~	0
2	6	423	rmse	standard	3.26	5	0.416	Preprocessor1_Model~	0
3	4	500	rmse	standard	3.28	5	0.442	Iter4	4
4	5	336	rmse	standard	3.29	5	0.446	Iter1	1
5	6	347	rmse	standard	3.29	5	0.411	Preprocessor1_Model~	0
6	6	500	rmse	standard	3.29	5	0.413	Iter6	6
7	7	500	rmse	standard	3.30	5	0.411	Iter3	3
8	8	399	rmse	standard	3.32	5	0.393	Preprocessor1_Model~	0
9	9	186	rmse	standard	3.33	5	0.367	Preprocessor1_Model~	0
10	9	477	rmse	standard	3.34	5	0.384	Preprocessor1_Model~	0



Here we try the default exponential correlation.

```

bayesres2<- tune_bayes(tune_wf,
  resamples = vfold,
  #metrics = rmse,
  #corr=list(type="matern",nu=5/2),
  #default in corr_mat(GPfit) is "exponential" power 1.95
  initial = 10,
  param_info = tune_param,
  iter = 10,
  objective=exp_improve(),
  control = ctrl
)
dput(bayesres2,"bayesres2.dd")

bayesres2=dget("bayesres2.dd")
show_best(bayesres2,n=10)

```


A tibble: 10 x 9

	mtry	trees	.metric	.estimator	mean	n	std_err	.config	.iter
	<int>	<int>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>	<int>
1	5	499	rmse	standard	3.25	5	0.447	Iter6	6
2	4	313	rmse	standard	3.29	5	0.452	Iter5	5
3	7	445	rmse	standard	3.29	5	0.399	Preprocessor1_Model~	0
4	5	453	rmse	standard	3.30	5	0.436	Iter8	8
5	6	498	rmse	standard	3.31	5	0.429	Iter2	2
6	4	500	rmse	standard	3.31	5	0.451	Iter1	1
7	7	500	rmse	standard	3.31	5	0.402	Iter9	9
8	5	258	rmse	standard	3.32	5	0.416	Preprocessor1_Model~	0
9	8	496	rmse	standard	3.33	5	0.398	Iter4	4
10	4	231	rmse	standard	3.33	5	0.433	Iter10	10

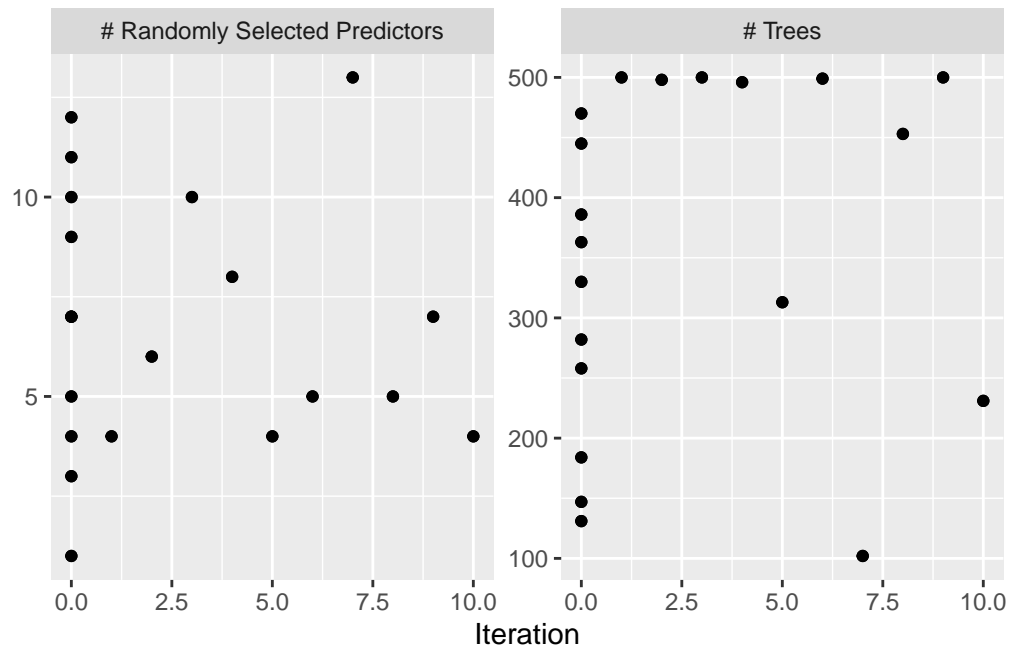
```

autoplot(bayesres2,type="performance")

```



```
autoplot(bayesres2,type="parameters")
```



2.3.2 Suggested software

(Frazier 2018, Ch 6)

- R: DiceOptim (on CRAN)
 - R: tune_bayes in `tune` (also CRAN)
 - Python: Spearmint <https://github.com/HIPS/Spearmint>
 - Python: GPyOpt <https://github.com/SheffieldML/GPyOpt>
 - Python: GPFlow (Tensorflow) <https://github.com/GPflow/GPflow> and GPyTorch (PyTorch) <https://github.com/cornellius-gp/gpytorch>
-

2.4 Design of experiments and response surface methodology

G. A. Lujan-Moreno, P. R. Howard, O. G. Rojas and D. C. Montgomery (2018): Design of experiments and response surface methodology to tune machine learning hyperparameters, with a random forest case- study. *Expert Systems with Applications*. 109, 195-205.

See separate slide-deck made by Håkon Gryvill, Yngvild Hamre and Javier Aguilar for the article presentation in MA8701 in the spring of 2021.

3 References

- M. Feurer and F. Hutter (2019). In F. Hutter et al (eds.) *Automated Machine Learning*. The Springer Series on Challenges in Machine Learning.
- Jo Eidsvik (2017): Introduction to Gaussian processes, note to TMA4265.
- Peter I. Frazier (2018): A tutorial on Bayesian optimization. arxiv <https://arxiv.org/abs/1807.02811>
- Max Kuhn and Julia Silge Version 0.0.1.9008 (2021-02-15) Tidy modelling with R. <https://www.tmwr.org/>
- Roger Gosse, University of Toronto: CSC321 Lecture 21: Bayesian Hyperparameter Optimization. http://www.cs.toronto.edu/~rgrosse/courses/csc321_2017/slides/lec21.pdf

- Max Kuhn (2020). caret: Classification and Regression Training. R package version 6.0-86. <https://CRAN.R-project.org/package=caret>